

Importing the libraries

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
import pandas as pd
```

Data Preparation

```
df = pd.read_csv("/content/drive/MyDrive/DV/transactions.csv")
```

```
df.head()
```

	Time	V1	V2	V3	V4	V5	V6	
V7 \								
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
0.239599								
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
0.078803								
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
0.791461								
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
0.237609								
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	
0.592941								
		V8	V9	...	V21	V22	V23	V24
V25 \								
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	
0.128539								
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	
0.167170								
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-
0.327642								
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	
0.647376								
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-
0.206010								
	V26	V27	V28	Amount	Class			
0	-0.189115	0.133558	-0.021053	149.62	0			
1	0.125895	-0.008983	0.014724	2.69	0			
2	-0.139097	-0.055353	-0.059752	378.66	0			
3	-0.221929	0.062723	0.061458	123.50	0			

```
4  0.502292  0.219422  0.215153  69.99      0
```

```
[5 rows x 31 columns]
```

```
df.shape
```

```
(284807, 31)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 284807 entries, 0 to 284806
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	Time	284807 non-null	float64
1	V1	284807 non-null	float64
2	V2	284807 non-null	float64
3	V3	284807 non-null	float64
4	V4	284807 non-null	float64
5	V5	284807 non-null	float64
6	V6	284807 non-null	float64
7	V7	284807 non-null	float64
8	V8	284807 non-null	float64
9	V9	284807 non-null	float64
10	V10	284807 non-null	float64
11	V11	284807 non-null	float64
12	V12	284807 non-null	float64
13	V13	284807 non-null	float64
14	V14	284807 non-null	float64
15	V15	284807 non-null	float64
16	V16	284807 non-null	float64
17	V17	284807 non-null	float64
18	V18	284807 non-null	float64
19	V19	284807 non-null	float64
20	V20	284807 non-null	float64
21	V21	284807 non-null	float64
22	V22	284807 non-null	float64
23	V23	284807 non-null	float64
24	V24	284807 non-null	float64
25	V25	284807 non-null	float64
26	V26	284807 non-null	float64
27	V27	284807 non-null	float64
28	V28	284807 non-null	float64
29	Amount	284807 non-null	float64
30	Class	284807 non-null	int64

```
dtypes: float64(30), int64(1)
```

```
memory usage: 67.4 MB
```

```
df.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```


#Frequency Distribution of Class

```
import matplotlib.pyplot as plt

class_counts = df['Class'].value_counts()

#Bar chart
plt.bar(class_counts.index, class_counts.values, color=['green',
'red'])
plt.xlabel('Class')
plt.ylabel('Transaction count')
plt.title('Class Distribution')
plt.xticks(class_counts.index, ['Non-Fraud', 'Fraud'])
# Adding annotations to each bar
for i, acc in enumerate(class_counts):
```

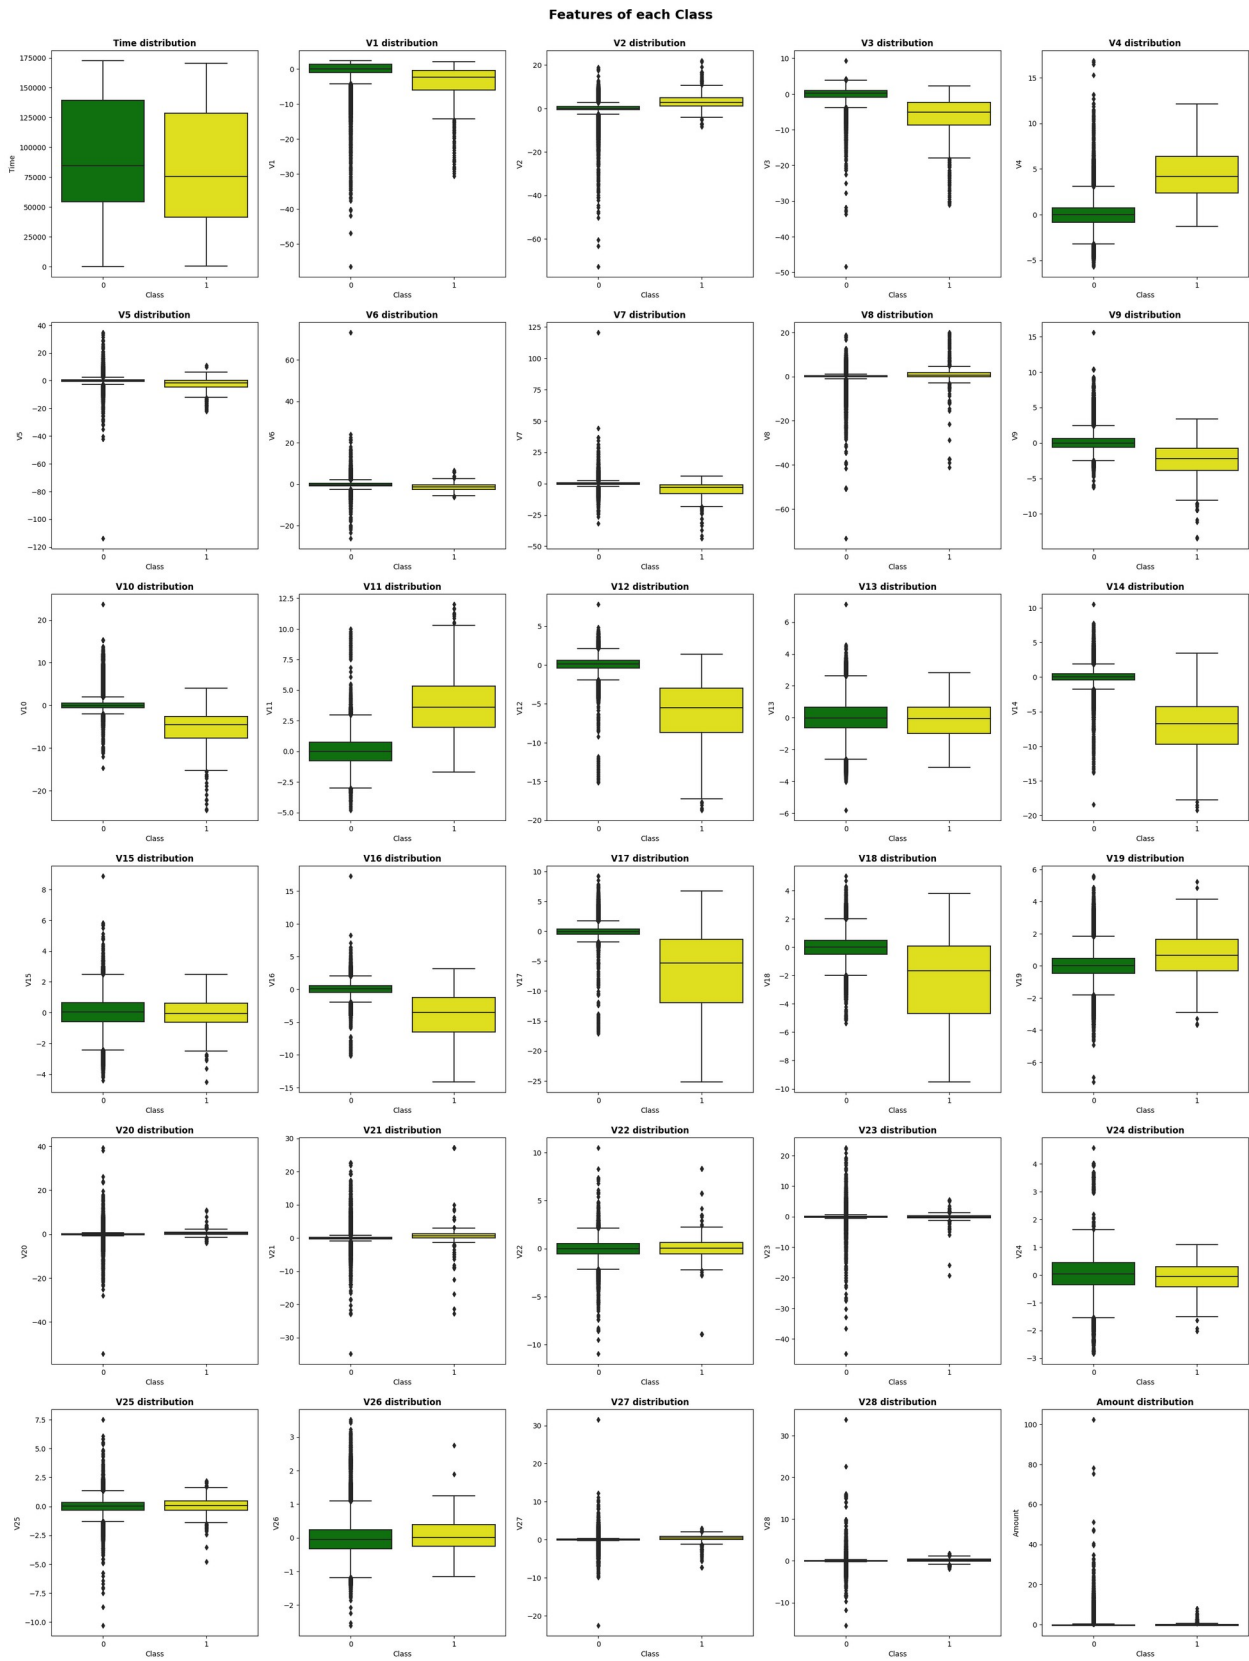
```
plt.text(i, acc + 0.5, acc, ha='center', va='bottom')
plt.show()
```



#Box Plot of Features of each Class

```
fig, axes = plt.subplots(nrows=6, ncols=5, figsize=(30,40))
fig.suptitle('Features of each Class', size = 18, y=0.9,
fontweight='bold')
fila=0
column=0

for i in df.columns[:-1]:
    sns.boxplot(ax=axes[fila,column], data=df, x='Class', y=i,
palette=['green','yellow'])
    axes[fila,column].set_title(f"{i} distribution",fontweight='bold')
    column= column+1
    if column == 5:
        column=0
        fila+=1
```




#Correlation Heat Map


```
# Creating subplots
fig, ax = plt.subplots(1, 2, figsize=(18, 6))

# Histogram for 'Time' feature
sns.histplot(df['Time'], bins=50, ax=ax[0], kde=True)
ax[0].set_title('Distribution of Transaction Time')
ax[0].set_xlabel('Time (seconds)')
ax[0].set_ylabel('Frequency')

# Histogram for 'Amount' feature
sns.histplot(df['Amount'], bins=50, ax=ax[1], kde=True)
ax[1].set_title('Distribution of Transaction Amount')
ax[1].set_xlabel('Amount')
ax[1].set_ylabel('Frequency')

# Adjust layout
plt.tight_layout()
plt.show()
```



The histograms for the 'Time' and 'Amount' features show the following:

1. **Transaction Time:** The distribution of transaction times shows some periodic patterns, which could be related to daily or weekly cycles in transaction frequency. This is common in financial transaction data, where activity can vary significantly depending on the time of day or week.
2. **Transaction Amount:** Most transactions are of lower amounts, with the frequency rapidly decreasing as the amount increases. This long-tailed distribution is typical in financial datasets, where small transactions are common and large transactions are relatively rare.

Feature Scaling

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
df['Amount'] = sc.fit_transform(pd.DataFrame(df['Amount']))
df.head()
```

	Time	V1	V2	V3	V4	V5	V6	
V7 \								
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
0.239599								
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
0.078803								
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
0.791461								
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
0.237609								
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	
0.592941								
		V8	V9	...	V21	V22	V23	V24
V25 \								
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	
0.128539								
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	
0.167170								
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	
0.327642								
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	
0.647376								
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	
0.206010								
	V26	V27	V28	Amount	Class			
0	-0.189115	0.133558	-0.021053	0.244964	0			
1	0.125895	-0.008983	0.014724	-0.342475	0			
2	-0.139097	-0.055353	-0.059752	1.160686	0			
3	-0.221929	0.062723	0.061458	0.140534	0			
4	0.502292	0.219422	0.215153	-0.073403	0			

```
[5 rows x 31 columns]
```

```
df['Amount'].describe()
```

```
count    2.848070e+05
mean     -1.596686e-17
std       1.000002e+00
min      -3.532294e-01
25%      -3.308401e-01
50%      -2.652715e-01
75%      -4.471707e-02
max       1.023622e+02
Name: Amount, dtype: float64
```



```
df = df.drop(['Time'],axis=1)
```

```
df.head()
```

	V1	V2	V3	V4	V5	V6
V7 \						
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
0.239599						
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
0.078803						
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
0.791461						
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
0.237609						
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921
0.592941						

	V8	V9	V10	...	V21	V22	V23
V24 \							
0	0.098698	0.363787	0.090794	...	-0.018307	0.277838	-0.110474
0.066928							
1	0.085102	-0.255425	-0.166974	...	-0.225775	-0.638672	0.101288
0.339846							
2	0.247676	-1.514654	0.207643	...	0.247998	0.771679	0.909412
0.689281							
3	0.377436	-1.387024	-0.054952	...	-0.108300	0.005274	-0.190321
1.175575							
4	-0.270533	0.817739	0.753074	...	-0.009431	0.798278	-0.137458
0.141267							

	V25	V26	V27	V28	Amount	Class
0	0.128539	-0.189115	0.133558	-0.021053	0.244964	0
1	0.167170	0.125895	-0.008983	0.014724	-0.342475	0
2	-0.327642	-0.139097	-0.055353	-0.059752	1.160686	0
3	0.647376	-0.221929	0.062723	0.061458	0.140534	0
4	-0.206010	0.502292	0.219422	0.215153	-0.073403	0

```
[5 rows x 30 columns]
```

```
df.shape
```

```
(284807, 30)
```

```
df.duplicated().any()
```

```
True
```

REMOVE DUPLICATED VALUES

```
df = df.drop_duplicates()
```

```
df.shape
```

```
(275663, 30)
```

```
284807 - 275663
```

```
9144
```

NOT HANDLING IMBALANCED

```
df['Class'].value_counts()
```

```
0    275190
```

```
1         473
```

```
Name: Class, dtype: int64
```

STORE FEATURE MATRIX IN x AND RESPONSE IN VECTOR y

```
X = df.drop('Class',axis=1)
```

```
y = df['Class']
```

SPLITTING THE DATASET INTO THE TRAINING SET AND TEST SET

```
from sklearn.model_selection import train_test_split
X_train,X_test , y_train, y_test =
train_test_split(X,y,test_size=0.20,random_state=42)
```

Handling Imbalanced Dataset

Undersampling

```
normal = df[df['Class']==0]
```

```
fraud = df[df['Class']==1]
```

```
normal.shape
```

```
(275190, 30)
```

```
fraud.shape
```

```
(473, 30)
```

```
normal_sample=normal.sample(n=473)
```

```
normal_sample.shape
```

```
(473, 30)
```

```
new_data = pd.concat([normal_sample,fraud],ignore_index=True)
```

```
new_data.head()
```

	V1	V2	V3	V4	V5	V6
V7 \						
0	-1.478278	1.872233	-1.248675	-1.602380	1.549607	-1.184753
2.337972						
1	2.092500	-1.006312	-1.835767	-0.672136	-0.568465	-1.470872
0.028210						
2	0.311762	-2.602832	-2.942923	1.880832	-0.090804	-0.480139
1.865720						
3	-0.625405	0.069174	-1.475244	-2.459043	2.304930	3.373235
0.033103						
4	2.095020	0.040510	-1.847722	0.338714	0.574470	-0.840730
0.485067						

	V8	V9	V10	...	V21	V22	V23
V24 \							
0	-1.049094	1.282513	2.288424	...	-0.168881	0.852646	-0.298441
0.657866							
1	-0.383493	-0.243969	0.836672	...	0.240261	0.550606	-0.039937
0.011346							
2	-0.549267	-0.303862	0.069705	...	0.718151	0.220383	-0.913711
0.332605							
3	1.197842	-1.841251	-0.157792	...	-0.348845	-0.877869	0.167975
0.682458							
4	-0.340795	0.577797	-0.077917	...	-0.013551	0.188564	0.062087
0.630235							

	V25	V26	V27	V28	Amount	Class
0	0.029204	0.011506	0.624286	0.194931	-0.199223	0
1	0.243252	0.006025	-0.083359	-0.069236	-0.013952	0
2	0.278026	-0.474820	-0.200782	0.073386	3.233054	0
3	-0.342840	0.450133	-0.034917	0.067017	-0.085478	0
4	0.386115	-0.243627	-0.036975	-0.061877	-0.349231	0

```
[5 rows x 30 columns]
```

```
new_data.shape
```

```
(946, 30)
```

```
new_data['Class'].value_counts()
```

```
0    473
1    473
Name: Class, dtype: int64
```


```

import matplotlib.pyplot as plt

class_counts = new_data['Class'].value_counts()

#Bar chart
plt.bar(class_counts.index, class_counts.values, color=['green',
'red'])
plt.xlabel('Class')
plt.ylabel('Transaction count')
plt.title('Class Distribution')
plt.xticks(class_counts.index, ['Non-Fraud', 'Fraud'])
# Adding annotations to each bar
for i, acc in enumerate(class_counts):
    plt.text(i, acc + 0.5, acc, ha='center', va='bottom')
plt.show()

```



```
new_data.head()
```

	V1	V2	V3	V4	V5	V6
0	1.900368	-0.754937	-0.029267	0.455184	-0.966945	0.332113
1	1.342034	-0.601087	-0.397937	-0.770990	-0.475795	-0.685272

```

0.072419
2  1.956950 -0.001275 -1.813899  0.603469  0.028823 -1.554520
0.277435
3  1.337227  0.290879 -0.387485  0.166632  0.567025  0.250590 -
0.025959
4  1.507577 -1.024466  0.559104 -1.413875 -1.543809 -0.704107 -
1.066982

      V8      V9      V10  ...      V21      V22      V23
V24  \
0  0.462351  1.537534  0.225682  ...  0.034383  0.074968  0.371245
0.640507
1 -0.243450 -0.945197  0.663080  ... -0.837065 -2.191765  0.112161 -
0.662780
2 -0.351745  0.737834 -0.507092  ...  0.209335  0.615825 -0.053269 -
0.052135
3  0.030460 -0.171999 -0.123140  ... -0.345309 -1.001003 -0.100966 -
1.427895
4 -0.159086 -1.726180  1.407877  ... -0.385481 -0.777162  0.125373 -
0.018079

      V25      V26      V27      V28      Amount      Class
0 -0.647078  0.133690 -0.012905 -0.042023 -0.293258      0
1  0.072125  0.680915 -0.095004  0.009330 -0.033502      0
2  0.198322 -0.105898 -0.017494 -0.022135 -0.136293      0
3  0.431179  0.182841 -0.026243  0.003852 -0.349831      0
4  0.177906 -0.389423  0.041005  0.028210 -0.258875      0

[5 rows x 30 columns]

X = new_data.drop('Class',axis=1)
y = new_data['Class']

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.20,random_state=42)

```

Logistic Regression

```

from sklearn.linear_model import LogisticRegression
log = LogisticRegression()
log.fit(X_train,y_train)

LogisticRegression()

y_pred1 = log.predict(X_test)

# Calculate evaluation metrics
from sklearn.metrics import accuracy_score, precision_score,

```

```

recall_score, f1_score
accuracy1 = accuracy_score(y_test, y_pred1)
precision = precision_score(y_test, y_pred1)
recall = recall_score(y_test, y_pred1)
f1 = f1_score(y_test, y_pred1)

metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy1, precision, recall, f1]
})
metrics_df

```

	Metric	Score
0	Accuracy	0.947368
1	Precision	0.989362
2	Recall	0.911765
3	F1 Score	0.948980

```

from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred1)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

```

```


sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

```

```

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()

```



		Predicted	
		Negative (N) -	Positive (P) +
Actual	Negative -	True Negative (TN)	False Positive (FP) Type I Error
	Positive +	False Negative (FN) Type II Error	True Positive (TP)

Decision Tree Classifier

```
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train,y_train)

DecisionTreeClassifier()

y_pred2 = dt.predict(X_test)

# Calculate evaluation metrics
accuracy2 = accuracy_score(y_test, y_pred2)
precision = precision_score(y_test, y_pred2)
recall = recall_score(y_test, y_pred2)
f1 = f1_score(y_test, y_pred2)

metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy2, precision, recall, f1]
})


print(metrics_df)
```

	Metric	Score
0	Accuracy	0.900000
1	Precision	0.902913
2	Recall	0.911765
3	F1 Score	0.907317

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred2)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()
```

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train,y_train)

RandomForestClassifier()

y_pred3 = rf.predict(X_test)

# Calculate evaluation metrics
accuracy3 = accuracy_score(y_test, y_pred3)
precision = precision_score(y_test, y_pred3)
recall = recall_score(y_test, y_pred3)
f1 = f1_score(y_test, y_pred3)

metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy3, precision, recall, f1]
})


print(metrics_df)
```

	Metric	Score
0	Accuracy	0.936842
1	Precision	0.978723
2	Recall	0.901961
3	F1 Score	0.938776

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred3)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]
```

```
sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)
```

```
plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()
```



Artificial Neural Network

```
import tensorflow as tf
from tensorflow import keras

# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build the ANN model
model = keras.Sequential([
    keras.layers.Dense(units=128, activation='relu',
input_dim=X_train.shape[1]),
    keras.layers.Dense(units=64, activation='relu'),
    keras.layers.Dense(units=1, activation='sigmoid') # Binary
classification, so using sigmoid activation
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

# Train the model
model.fit(X_train_scaled, y_train, epochs=10, batch_size=32,
validation_split=0.2)

Epoch 1/10
19/19 [=====] - 2s 16ms/step - loss: 0.4642 -
accuracy: 0.8477 - val_loss: 0.3449 - val_accuracy: 0.8882
Epoch 2/10
19/19 [=====] - 0s 6ms/step - loss: 0.2727 -
accuracy: 0.9089 - val_loss: 0.2432 - val_accuracy: 0.9079
Epoch 3/10
19/19 [=====] - 0s 7ms/step - loss: 0.2114 -
accuracy: 0.9305 - val_loss: 0.1891 - val_accuracy: 0.9145
Epoch 4/10
19/19 [=====] - 0s 8ms/step - loss: 0.1851 -
accuracy: 0.9371 - val_loss: 0.1682 - val_accuracy: 0.9211
Epoch 5/10
19/19 [=====] - 0s 10ms/step - loss: 0.1699 -
accuracy: 0.9437 - val_loss: 0.1462 - val_accuracy: 0.9539
Epoch 6/10
19/19 [=====] - 0s 7ms/step - loss: 0.1563 -
accuracy: 0.9470 - val_loss: 0.1360 - val_accuracy: 0.9474
Epoch 7/10
19/19 [=====] - 0s 5ms/step - loss: 0.1446 -
accuracy: 0.9487 - val_loss: 0.1294 - val_accuracy: 0.9539
Epoch 8/10
19/19 [=====] - 0s 7ms/step - loss: 0.1347 -
```

```

accuracy: 0.9503 - val_loss: 0.1212 - val_accuracy: 0.9671
Epoch 9/10
19/19 [=====] - 0s 6ms/step - loss: 0.1270 -
accuracy: 0.9503 - val_loss: 0.1200 - val_accuracy: 0.9539
Epoch 10/10
19/19 [=====] - 0s 6ms/step - loss: 0.1217 -
accuracy: 0.9570 - val_loss: 0.1185 - val_accuracy: 0.9605

<keras.src.callbacks.History at 0x7bb58c16bf40>

# Make predictions on the test set
y_pred_proba = model.predict(X_test_scaled)
y_pred4 = (y_pred_proba > 0.5).astype(int)

6/6 [=====] - 0s 2ms/step

# Evaluate the model
accuracy4 = accuracy_score(y_test, y_pred4)
accuracy4

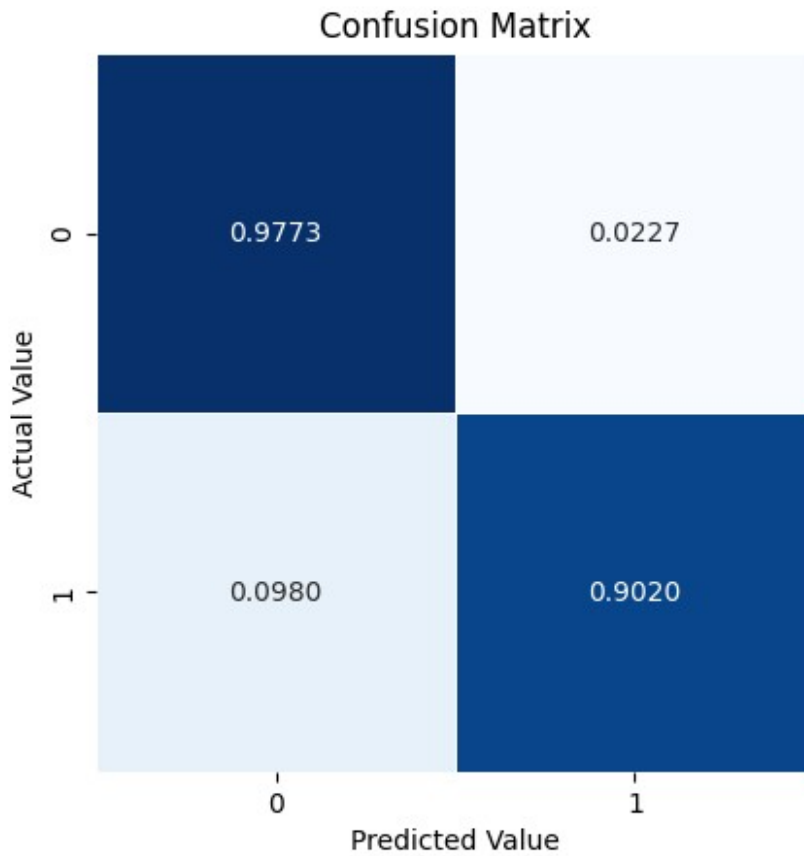
0.9368421052631579

from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred4)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()

```



#Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier

# Train the Gradient Boosting classifier
gb_classifier = GradientBoostingClassifier()
gb_classifier.fit(X_train, y_train)

GradientBoostingClassifier()

# Make predictions on the test data
y_pred5 = gb_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy5 = accuracy_score(y_test, y_pred5)
precision_gb = precision_score(y_test, y_pred5)
recall_gb = recall_score(y_test, y_pred5)
f1_gb = f1_score(y_test, y_pred5)

# Display the metrics
metrics_gb_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy5, precision_gb, recall_gb, f1_gb]
})
```


```
print(metrics_gb_df)
```

	Metric	Score
0	Accuracy	0.931579
1	Precision	0.968421
2	Recall	0.901961
3	F1 Score	0.934010

```
from sklearn.metrics import confusion_matrix  
conf_matrix = confusion_matrix(y_test, y_pred5)  
conf_matrix_norm = conf_matrix.astype('float') /  
conf_matrix.sum(axis=1)[:, np.newaxis]
```

```
sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',  
cbar=False, linewidths=.5, square=True)
```

```
plt.xlabel('Predicted Value')  
plt.ylabel('Actual Value')  
plt.title('Confusion Matrix')  
plt.show()
```



Results

```
final_data = pd.DataFrame({'Models': ['LR', 'DT', 'RF', 'ANN', 'GB'],
                           "ACC": [accuracy1*100,
                                   accuracy2*100,
                                   accuracy3*100,
                                   accuracy4*100,
                                   accuracy5*100,
                                   ]})
```

final_data

	Models	ACC
0	LR	94.736842
1	DT	90.000000
2	RF	93.684211
3	ANN	93.684211
4	GB	93.157895

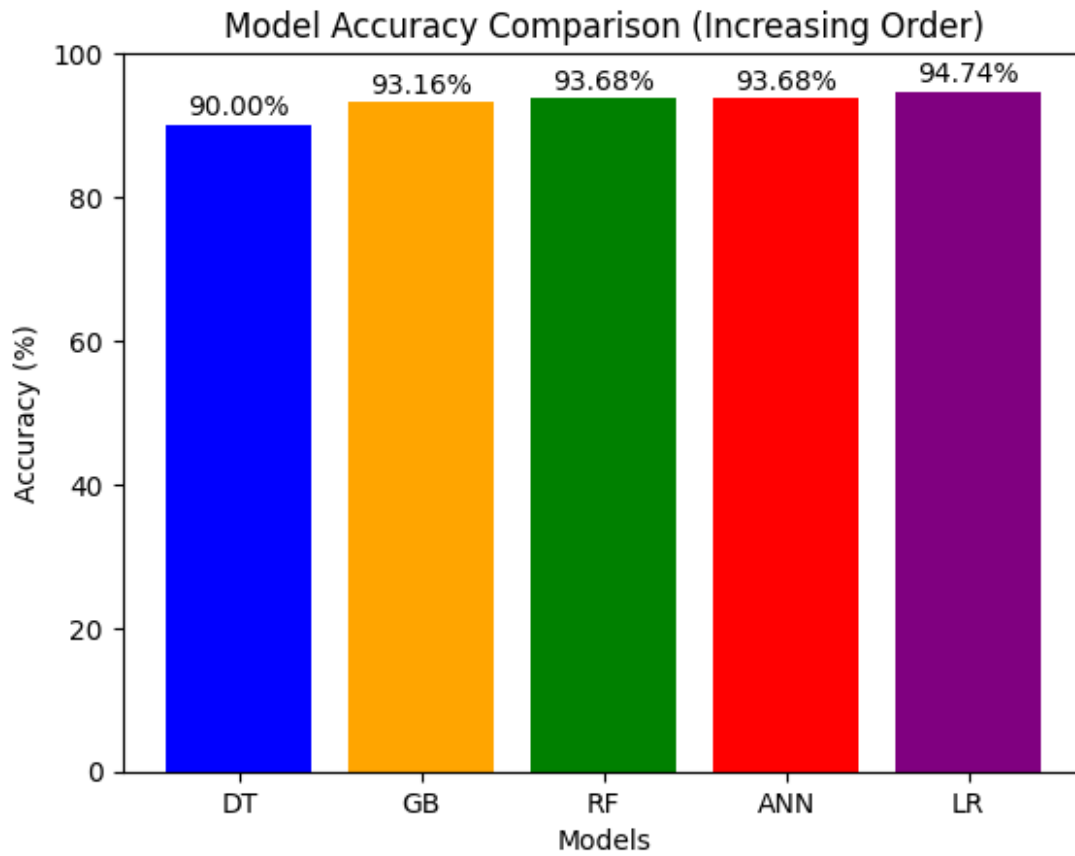
```
import matplotlib.pyplot as plt

# Sort the DataFrame by the 'ACC' column in ascending order
final_data_sorted = final_data.sort_values(by='ACC')

# Plotting the bar chart
plt.bar(final_data_sorted['Models'], final_data_sorted['ACC'],
        color=['blue', 'orange', 'green', 'red', 'purple'])
plt.xlabel('Models')
plt.ylabel('Accuracy (%)')
plt.title('Model Accuracy Comparison (Increasing Order)')
plt.ylim(0, 100) # Set the y-axis limits to ensure proper
visualization of percentages

# Adding annotations to each bar
for i, acc in enumerate(final_data_sorted['ACC']):
    plt.text(i, acc + 0.5, f'{acc:.2f}%', ha='center', va='bottom')

plt.show()
```



Oversampling

```
X = df.drop('Class', axis=1)
y=df['Class']
X.shape
(275663, 29)
y.shape
(275663,)
from imblearn.over_sampling import SMOTE
X_res,y_res = SMOTE().fit_resample(X,y)
y_res.value_counts()
0    275190
1    275190
Name: Class, dtype: int64
```




```

import matplotlib.pyplot as plt

class_counts = y_res.value_counts()

#Bar chart
plt.bar(class_counts.index, class_counts.values, color=['green',
'red'])
plt.xlabel('Class')
plt.ylabel('Transaction count')
plt.title('Class Distribution')
plt.xticks(class_counts.index, ['Non-Fraud', 'Fraud'])
# Adding annotations to each bar
for i, acc in enumerate(class_counts):
    plt.text(i, acc + 0.5, acc, ha='center', va='bottom')
plt.show()

```



```

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X_res,y_res,test_size=0.20,random_state=42)

```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
log = LogisticRegression()
log.fit(X_train,y_train)

LogisticRegression()

y_pred1 = log.predict(X_test)

# Calculate evaluation metrics
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
accuracy1 = accuracy_score(y_test, y_pred1)
precision = precision_score(y_test, y_pred1)
recall = recall_score(y_test, y_pred1)
f1 = f1_score(y_test, y_pred1)


metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy1, precision, recall, f1]
})
metrics_df
```

	Metric	Score
0	Accuracy	0.944148
1	Precision	0.973062
2	Recall	0.913514
3	F1 Score	0.942348

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred1)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()
```



Decision Tree Classifier

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier()
dt.fit(X_train,y_train)

DecisionTreeClassifier()

y_pred2 = dt.predict(X_test)

# Calculate evaluation metrics
accuracy2 = accuracy_score(y_test, y_pred2)
precision = precision_score(y_test, y_pred2)
recall = recall_score(y_test, y_pred2)
f1 = f1_score(y_test, y_pred2)

metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy2, precision, recall, f1]
})


print(metrics_df)
```

	Metric	Score
0	Accuracy	0.997892
1	Precision	0.996915
2	Recall	0.998873
3	F1 Score	0.997893

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred2)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]
```

```
sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)
```

```
plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()
```



Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
```

```
RandomForestClassifier()
```

```
y_pred3 = rf.predict(X_test)
```

```
# Calculate evaluation metrics
```

```
accuracy3 = accuracy_score(y_test, y_pred3)
precision = precision_score(y_test, y_pred3)
recall = recall_score(y_test, y_pred3)
f1 = f1_score(y_test, y_pred3)
```

```
metrics_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy3, precision, recall, f1]
})
```


```
print(metrics_df)
```

	Metric	Score
0	Accuracy	0.999927
1	Precision	0.999855
2	Recall	1.000000
3	F1 Score	0.999927

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred3)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]
```

```
sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)
```

```
plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()
```



Artificial Neural Network

```
import tensorflow as tf
from tensorflow import keras

# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build the ANN model
model = keras.Sequential([
    keras.layers.Dense(units=128, activation='relu',
input_dim=X_train.shape[1]),
    keras.layers.Dense(units=64, activation='relu'),
    keras.layers.Dense(units=1, activation='sigmoid') # Binary
classification, so using sigmoid activation
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
```

```
# Train the model
```

```
model.fit(X_train_scaled, y_train, epochs=10, batch_size=32,  
validation_split=0.2)
```

```
Epoch 1/10
```

```
11008/11008 [=====] - 31s 3ms/step - loss:  
0.0285 - accuracy: 0.9906 - val_loss: 0.0077 - val_accuracy: 0.9981
```

```
Epoch 2/10
```

```
11008/11008 [=====] - 30s 3ms/step - loss:  
0.0078 - accuracy: 0.9979 - val_loss: 0.0052 - val_accuracy: 0.9986
```

```
Epoch 3/10
```

```
11008/11008 [=====] - 29s 3ms/step - loss:  
0.0055 - accuracy: 0.9985 - val_loss: 0.0077 - val_accuracy: 0.9976
```

```
Epoch 4/10
```

```
11008/11008 [=====] - 31s 3ms/step - loss:  
0.0042 - accuracy: 0.9988 - val_loss: 0.0037 - val_accuracy: 0.9991
```

```
Epoch 5/10
```

```
11008/11008 [=====] - 31s 3ms/step - loss:  
0.0034 - accuracy: 0.9991 - val_loss: 0.0053 - val_accuracy: 0.9988
```

```
Epoch 6/10
```

```
11008/11008 [=====] - 31s 3ms/step - loss:  
0.0033 - accuracy: 0.9991 - val_loss: 0.0026 - val_accuracy: 0.9993
```

```
Epoch 7/10
```

```
11008/11008 [=====] - 48s 4ms/step - loss:  
0.0027 - accuracy: 0.9992 - val_loss: 0.0038 - val_accuracy: 0.9994
```

```
Epoch 8/10
```

```
11008/11008 [=====] - 28s 3ms/step - loss:  
0.0026 - accuracy: 0.9994 - val_loss: 0.0032 - val_accuracy: 0.9992
```

```
Epoch 9/10
```

```
11008/11008 [=====] - 30s 3ms/step - loss:  
0.0024 - accuracy: 0.9994 - val_loss: 0.0029 - val_accuracy: 0.9993
```

```
Epoch 10/10
```

```
11008/11008 [=====] - 30s 3ms/step - loss:  
0.0021 - accuracy: 0.9995 - val_loss: 0.0020 - val_accuracy: 0.9995
```

```
<keras.src.callbacks.History at 0x7bb539fd4f70>
```

```
# Make predictions on the test set
```

```
y_pred_proba = model.predict(X_test_scaled)
```

```
y_pred4 = (y_pred_proba > 0.5).astype(int)
```

```
3440/3440 [=====] - 5s 1ms/step
```

```
# Evaluate the model
```

```
accuracy4 = accuracy_score(y_test, y_pred4)
```

```
accuracy4
```

```
0.9993731603619318
```

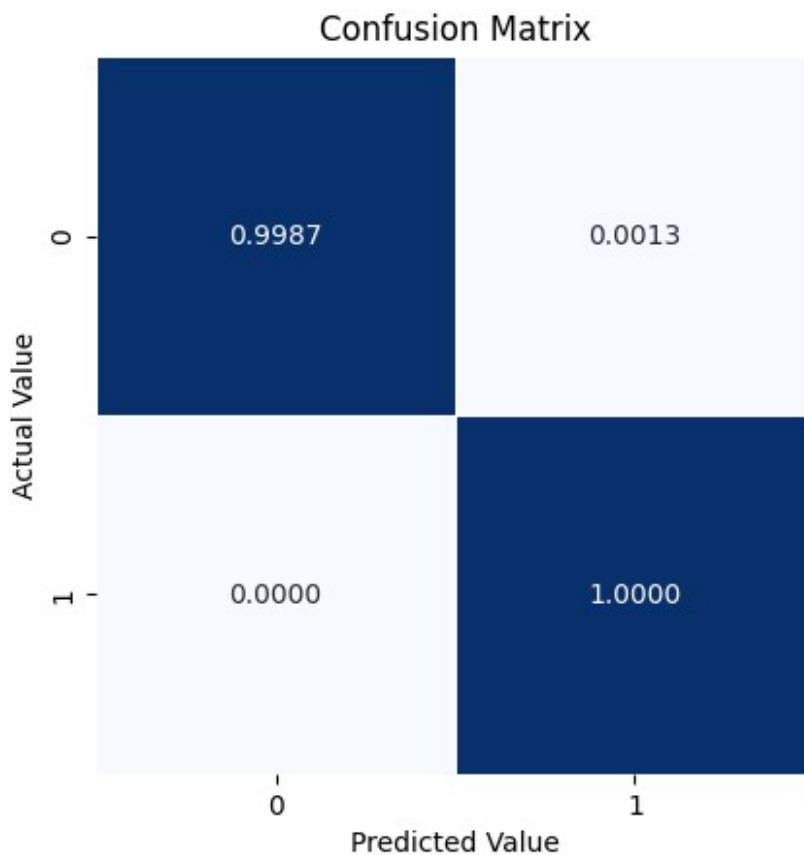
```

from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred4)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()

```



#Gradient Boosting

```

from sklearn.ensemble import GradientBoostingClassifier

# Train the Gradient Boosting classifier
gb_classifier = GradientBoostingClassifier()
gb_classifier.fit(X_train, y_train)

GradientBoostingClassifier()

```



```

# Make predictions on the test data
y_pred5 = gb_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy5 = accuracy_score(y_test, y_pred5)
precision_gb = precision_score(y_test, y_pred5)
recall_gb = recall_score(y_test, y_pred5)
f1_gb = f1_score(y_test, y_pred5)

# Display the metrics
metrics_gb_df = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Score': [accuracy5, precision_gb, recall_gb, f1_gb]
})

```

```
print(metrics_gb_df)
```

	Metric	Score
0	Accuracy	0.976353
1	Precision	0.986266
2	Recall	0.966129
3	F1 Score	0.976094

```

from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(y_test, y_pred5)
conf_matrix_norm = conf_matrix.astype('float') /
conf_matrix.sum(axis=1)[:, np.newaxis]

```

```


sns.heatmap(conf_matrix_norm, annot=True, fmt='.4f', cmap='Blues',
cbar=False, linewidths=.5, square=True)

```

```

plt.xlabel('Predicted Value')
plt.ylabel('Actual Value')
plt.title('Confusion Matrix')
plt.show()

```



#Results

```
final_data = pd.DataFrame({'Models': ['LR', 'DT', 'RF', 'ANN', 'GB'],
                           "ACC": [accuracy1*100,
                                    accuracy2*100,
                                    accuracy3*100,
                                    accuracy4*100,
                                    accuracy5*100,
                                    ]})
```

final_data

	Models	ACC
0	LR	94.414768
1	DT	99.789237
2	RF	99.992732
3	ANN	99.937316
4	GB	97.635270

```
import matplotlib.pyplot as plt
```

```
# Sort the DataFrame by the 'ACC' column in ascending order
final_data_sorted = final_data.sort_values(by='ACC')
```

```
# Plotting the bar chart
plt.bar(final_data_sorted['Models'], final_data_sorted['ACC'],
color=['blue', 'orange', 'green', 'red', 'purple'])
plt.xlabel('Models')
plt.ylabel('Accuracy (%)')
plt.title('Model Accuracy Comparison (Increasing Order)')
plt.ylim(0, 110) # Set the y-axis limits to ensure proper
visualization of percentages

# Adding annotations to each bar
for i, acc in enumerate(final_data_sorted['ACC']):
    plt.text(i, acc + 0.5, f'{acc:.2f}%', ha='center', va='bottom')

plt.show()
```

